

Linear Containers

Multidimensional Arrays,
Lists, Queues, Stacks



Georgi Georgiev
A guy that knows C++


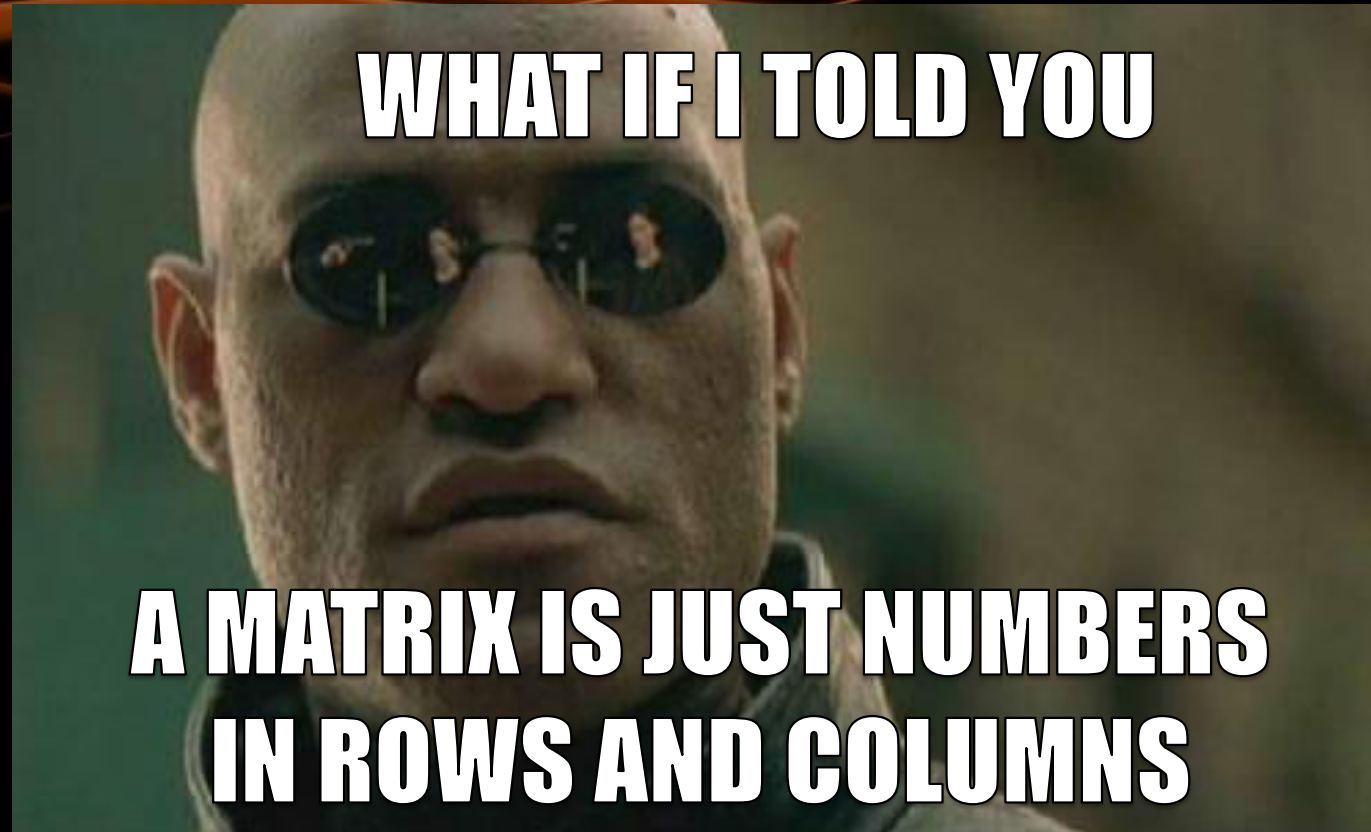


Table of Contents

1. Multidimensional Arrays
2. Data Structures - Concept
3. STL Linear Containers
 - **std::vector**
 - Iterators
 - **std::list**
 - **std::stack & std::queue**





Multidimensional Arrays

Matrices and Higher Dimensions

Multidimensional Arrays

- C++ can make arrays act "as if" they have many dimensions
 - "as if" – they are just normal arrays which are indexed differently
 - Compiler enforces dimension syntax in code
- Imagine each element is actually an array
 - 2D (matrix): array of arrays (each element is a "normal" array)
 - 3D array: array of 2D arrays (each element is a matrix)
- Most-common usage: making a matrix/table

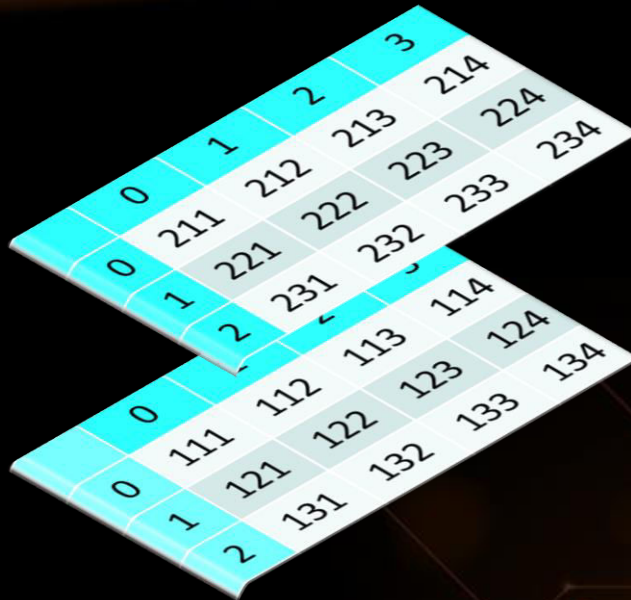
Using Multidimensional Arrays

- Accessing elements is done with one indexer per dimension
 - **matrix** 1st element (column) of 2nd row is **matrix[1][0]**
- Declaring: add a **[size]** for each additional dimension
 - First dimension can omit size
 - E.g. **int matrix2Rows3Cols[2][3]**
or just **matrix2Rows3Cols[][3]**
(needs initializer, unless function parameter)

Using Multidimensional Arrays

- Each dimension is an array with 1 less dimension
 - `int matrix2Rows3Cols[][3] = {
 {11, 12, 13},
 {21, 22, 23}
};`
 - `int cube[2][3][4] = {
 { {111, 112, 113, 114}, {121, 122, 123, 124}, {131, 132, 133, 134} },
 { {211, 212, 213, 214}, {221, 222, 223, 224}, {231, 232, 233, 234} }
};`
- If no initializer `{ }` brackets, values are undefined
- If more elements than initialized, others are defaults

	0					1			
	0	1	2	3		0	1	2	3
0	111	112	113	114	0	211	212	213	214
1	121	122	123	124	1	221	222	223	224
2	131	132	133	134	2	231	232	233	234



0	1	2	3	
0	111	112	113	114
1	121	122	123	124
2	131	132	133	134

Multidimensional Arrays

LIVE DEMO

Quick Quiz

TIME:

- What will the following code do?
- a) cause a compile-time error
- b) cause a runtime error due to index being out of bounds
- c) set `matrix[2][0]` to 0
- d) summon demons
- e) you know nothing

```
const int rows = 4;  
const int cols = 3;  
int matrix[rows][cols] = {  
    {11, 12, 13},  
    {21, 22, 23},  
    {31, 32, 33},  
    {41, 42, 43}  
};  
  
matrix[1][3] = 0;
```


0,0	0,1	0,2
1,0	1,1	1,2
2,0	2,1	2,2

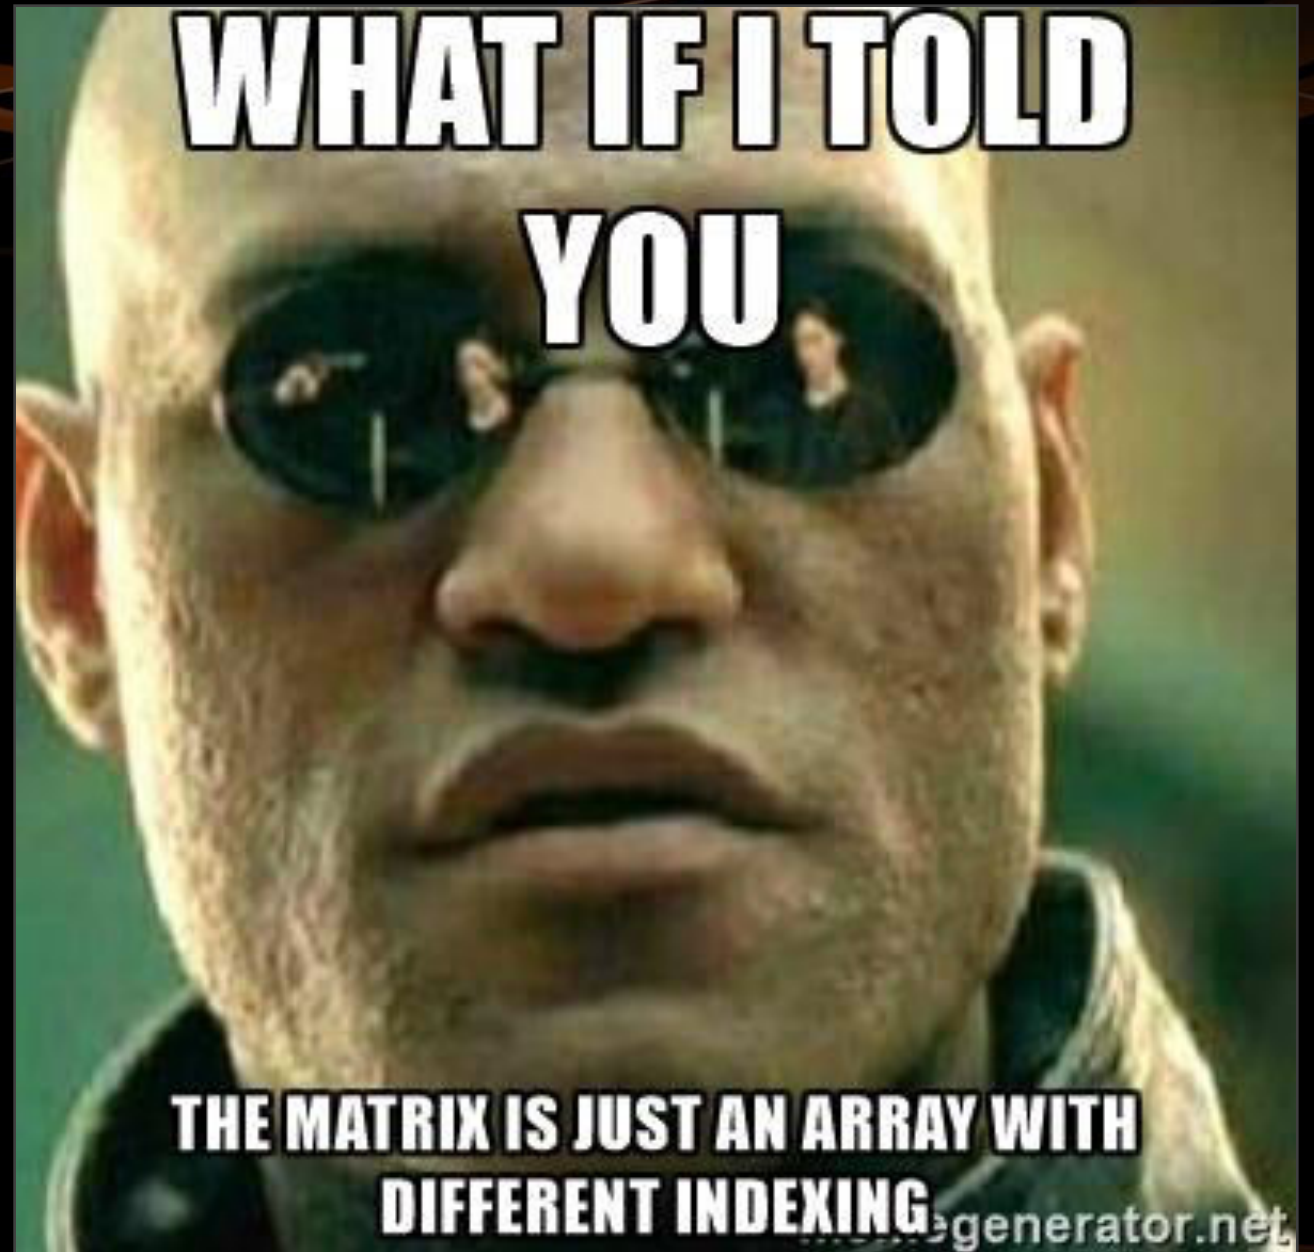
0,2	1,0	1,1	1,2	2,0
-----	-----	-----	-----	-----

C++ PITFALL: “OUT OF BOUNDS INSIDE” MULTIDIMENSIONAL ARRAYS

C++ (C actually) stores multidimensional arrays as 1D, by joining up together 1st dimension elements, e.g. for 2D arrays – joining up rows into a 1D array.

This is called “row-major order”

E.g. for a `matrix[rows][cols]` accessing `[r][c]` just means `[r * cols + c]` in the actual array



0,0	0,1	0,2
1,0	1,1	1,2
2,0	2,1	2,2

0,2	1,0	1,1	1,2	2,0
-----	-----	-----	-----	-----

Row-Major Order in Multidimensional Arrays

LIVE DEMO

"Multidimensional" Containers

- We know `std::vector` can contain any type
 - ... actually any type with a default constructor
 - `int`, `double`, `char`, `string`, even another `std::vector`, etc.
- Often containers (e.g. `vectors`) will contain other containers
- E.g. a vector of vectors (2D), a vector of vector of vectors (3D)
 - Element access is the same code as with multidimensional arrays
 - Note: no row-major order (not contiguous in memory)

"Multidimensional" Containers

LIVE DEMO

Data Structures

Classifying Data Containers by Operation

- Data Structures organize data for efficient access
 - Different data structures are efficient for different use-cases
 - Essentially – a data container + algorithms for access
- Some of the common data structures in computer science:
 - Arrays – fast access by index, constant/dynamic size
 - Linked-list – fast add/remove at any position, no index access
 - Map/Dictionary – contains key/value pairs, fast access by key

Complexity 101

- Complexity in Computer Science describes performance
 - How fast an algorithm runs & How much memory it consumes
 - Based on the size of the input data – usually denoted as N
 - We usually care about the worst-case performance
- How do we measure complexity?
 - time = number of basic steps, memory = number of elements
- Complexity is usually denoted by the Big-O notation
 - How much the number of steps grows compared to input size

Complexity 101

- We usually care about **X** orders of magnitude, not **+X** or ***X**
 - **$O(N+3)$** == **$O(2N)$** == **$O(N)$** , i.e. we care about the **N** parts
 - *If something takes 1 million or 2 million years, it's the "million" that bothers you, not the "1" or the "2"*
- **$O(1)$** – "constant" time/memory – input size has no effect
- **$O(\log(N))$** – logarithmic – complexity grows as $\log(\text{input})$ grows
- **$O(N)$** – linear – complexity grows as input grows
- **$O(N^2)$, $O(N^3)$, ...** – quadratic, cubic, ... – complexity grows with square/cube/etc. of input size
- **$O(2^N)$, $O(3^N)$, ...** – exponential – this is a monster

Data Structure Performance 101

- If **N** is the number of elements in the container (the **.size()**):

	vector	list	map, set	unordered_map, unordered_set
access i^{th}	$O(1)$	$O(i)$	$O(i)$	---
find(V)	$O(N)$	$O(N)$	$O(\log(N))$	$O(1)$
insert(V)	$O(1)$ at end (usually), $O(N)$ otherwise	$O(1)$	$O(\log(N))$	$O(1)$
erase(V)	$O(1)$ at end (usually), $O(N)$ otherwise	$O(1)$	$O(\log(N))$	$O(1)$
Getting a sorted sequence	$O(N \cdot \log(N))$ (using <code>std::sort</code> algorithm)	$O(N + N \cdot \log(N))$ (using <code>.sort()</code> method)	$O(N)$ (by just iterating)	---

STL Linear Containers

Vectors, Lists, Iterators, Container Adapters

std::vector

- Represents an array, has all array operations (i.e. **operator[]**)
- Changes size automatically when elements added
- **push_back()** complexity is *amortized* **$O(1)$**
 - i.e. usually takes **$O(1)$** time, occasionally takes **$O(N)$** time
 - i.e. slow ~10 times out of ~1000, ~32 times out of ~4 billion, etc.
- Fast access **$O(1)$** to any element (random index access)
 - **arr[0] = 69; arr[15] = 42;**

`std::vector`

LIVE DEMO

size_t and size_type

- Alias of one of the integer types
 - E.g. **unsigned long int** or **unsigned long long int**
 - Able to represent the size of any object in bytes
 - E.g. **sizeof()** returns **size_t**
- Each STL container offers a similar **::size_type**
 - A good practice is to use it instead of **int** for sizes, positions, etc.

```
for (vector<int>::size_type i = 0; i < nums.size(); i++) {  
    cout << nums[i] << endl;  
}
```

Container Iterators

- STL Iterators are things that know how to traverse a container
 - **operator++** - moves iterator to the next element
 - **operator*** - accesses the element
 - **operator->** - same as dot **operator.** on the element
- Each container has an iterator (e.g. **std::vector<T>::iterator**)
- Each container has **begin()** and **end()** iterators
 - **begin()** points to first element, **end()** to AFTER last
 - Range-based **for** loop uses those to work on ANY container

Using Iterators with Vectors

- Using iterators on **vectors** is almost the same as using indexes
- To walk over a vector:
 - Start from **begin()**, move with **++** until you reach **end()**
 - Access the current element with *****

```
vector<int> nums {42, 13, 69};  
for (vector<int>::iterator i = nums.begin(); i != nums.end(); i++) {  
    cout << *i << endl;  
}  
  
// Equivalent code  
for (vector<int>::size_type i = 0; i < nums.size(); i++) {  
    cout << nums[i] << endl;  
}
```


- Example: Change each element in the vector by dividing it by 2

```
vector<int> numbers {42, 13, 69};  
for (vector<int>::iterator i = numbers.begin(); i != numbers.end(); i++) {  
    *i /= 2;  
}  
// Equivalent code  
for (int i = 0; i < numbers.size(); i++) {  
    numbers[i] /= 2;  
}
```

- Example: Print each string element and its length

```
vector<string> words {"the", "quick", "purple", "fox"};  
for (vector<string>::iterator i = words.begin(); i != words.end(); i++) {  
    cout << *i << ": " << i->size() << endl;  
}  
// Equivalent code  
for (int i = 0; i < words.size(); i++) {  
    cout << words[i] << ": " << words[i].size() << endl;  
}
```

Using Iterators with Vectors

LIVE DEMO

Why Use Iterators?

- Vectors may not need iterators, because they have indexes
 - i.e. they have sequential elements accessible by **operator[]**
- Not all containers have indexes
 - Only **std::array**, **std::vector** & **std::deque** have indexes
 - The other containers don't offer access by index
- Iterators work on all containers, abstract-away container details
 - Doesn't matter what container you iterate, code is the same

std::list

- Represents elements connected to each other in a sequence
 - `std::list<int> values; std::list<string> names;`
 - Each element connects to the previous and next element.
Like Christmas lights
- All element access is done with iterators
- Can add or remove elements anywhere in **$O(1)$** time
 - Requires iterator to where an element should be added/removed
- `push_back()`, `push_front()`, `insert()`, `size()`, ...

`std::list`

LIVE DEMO

Container Adaptors

- Wrap a container (e.g. **vector**) with a different interface
- Allow you to express intentions better
- Make code more abstract and focused on the task, not the code
- STL Adapters for common Computer Science data structures:
 - **std::stack** is a first-in, last-out (FIFO) data structure
 - **std::queue** is a first-in, first-out (FIFO) data structure
 - **std::priority_queue** is a data structure that gives quick access to the "highest priority item"

std::stack

- Represents a container (a **deque** by default) working like a stack
- A stack is a "first-in, last-out structure" (FILO)
- *Imagine a pile of bricks*
 - *the last brick you put is the first you can remove*
- Access to elements other than **top()** is not provided
- **top()** gets the top, **pop()** removes it, **push(T)** adds to top

std::queue

- Represents a container (**deque** by default) working like a queue
- A queue is a "first-in, first-out" structure (FIFO)
- *Imagine a line at a store*
 - *first person in the line is the first to get out*
- Access to elements other than **front()** is not provided
- **front()** gets first, **pop()** removes it, **push(T)** adds to back

std::priority queue

- Represents a queue, but elements are ordered by priority
 - By default, "larger" elements have higher priority
- *Imagine a queue at a hospital's emergency room*
 - *Patients with more serious cases treated BEFORE those with less serious ones*
- Higher priority elements move in front of lower priority ones
 - Getting top-priority element is **$O(1)$** , insertion is **$\log(N)$**
 - To get top-priority element use **`top()`** (instead of **`front()`**)

Container Adaptors

LIVE DEMO

Summary

- C++ Multidimensional arrays are just normal arrays
 - ... with indexing for each dimension
- We usually measure performance based on input
 - We care how quickly much performance degrades based on input size. We use Big-O notation to denote that
- C++ offers several linear data structures and adaptors
 - **vector, list, stack, queue, priority_queue**
 - Each is efficient for certain use-cases



Linear Containers



Questions?

